

Name: Abdul Rafay
BS Software Engineering (4th)
Arch Technologies Category B – Task 2

Chapter 1: Machine Learning Landscape

➤ What is Machine Learning?

Machine Learning is all about teaching computers to learn from data and make decisions without being specifically programmed for every task. A helpful way to think about it is through a definition by Tom Mitchell:

A computer learns when it gets better at a task (T), using experience (E), and its performance improves over time, as measured by some criteria (P).

➤ Types of Machine Learning

- **Supervised Learning:** The model learns using labeled data, meaning we already know the answers. It's used in tasks like predicting prices (regression) or identifying objects (classification).
- **Unsupervised Learning:** Here, the data has no labels. The goal is to find patterns or groupings. Examples include clustering and dimensionality reduction (like PCA).
- **Semi-supervised Learning:** A mix of both labeled and unlabeled data. This is useful when labeling data is expensive or time-consuming.
- **Reinforcement Learning:** The model learns by interacting with an environment and receiving feedback, like rewards or penalties, much like training a pet.

➤ How Models Learn

- **Batch Learning:** The system is trained all at once using all the available data.
- **Online Learning:** The model learns little by little, as new data comes in. This is useful for real-time systems.

➤ Learning Techniques

- **Instance-Based Learning:** This method works by comparing new problems with past examples, kind of like looking up similar cases. An example is k-Nearest Neighbors.

- **Model-Based Learning:** In this case, the system tries to build a general model from the training data. Linear regression is a good example.

➤ Common Challenges

Some typical issues in machine learning include not having enough data, poor-quality data, overfitting (model learns too much from the training data and doesn't generalize), underfitting (model is too simple), and tuning hyperparameters to get better results.

Chapter 2: End to End Machine Learning Project

➤ Important Libraries:

```
# Importing essential libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

➤ Scikit-Learn Tools:

```
from sklearn.model_selection import train_test_split, cross_val_score, GridSearchCV
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import Pipeline
from sklearn.compose import ColumnTransformer
from sklearn.linear_model import LinearRegression
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error
from sklearn.datasets import fetch_california_housing
```

➤ Dataset:

```
housing = fetch_california_housing(as_frame=True)
df = housing.frame
```

➤ Initial Data Check:

```
print(df.head())
print(df.info())
print(df.describe())
```

➤ Histogram:

```
df.hist(bins=50, figsize=(20, 15))
plt.show()
```

➤ Correlation matrix heatmap:

```
corr_matrix = df.corr()
plt.figure(figsize=(12, 8))
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm')
plt.title("Correlation Heatmap")
plt.show()
```

➤ Scatterplot:

```
df.plot(kind="scatter", x="MedInc", y="MedHouseVal", alpha=0.1)
plt.title("Median Income vs. Median House Value")
plt.show()
```

➤ Splitting the data:

```
train_set, test_set = train_test_split(df, test_size=0.2, random_state=42)
housing_train = train_set.copy()
housing_test = test_set.copy()
```

➤ Separating features and labels:

```
housing = housing_train.drop("MedHouseVal", axis=1)
housing_labels = housing_train["MedHouseVal"].copy()
```

➤ Numerical Features:

```
num_attribs = housing.select_dtypes(include=[np.number]).columns.tolist()
```

➤ Pipeline for numerical data preprocessing:

```
num_pipeline = Pipeline([
    ('imputer', SimpleImputer(strategy="median")),
    ('std_scaler', StandardScaler()),
])
```

➤ Full preprocessing pipeline:

```
full_pipeline = ColumnTransformer([
    ("num", num_pipeline, num_attribs),
])
```

➤ Preparing data:

```
housing_prepared = full_pipeline.fit_transform(housing)
```

➤ Models trained:

```
lin_reg = LinearRegression()
lin_reg.fit(housing_prepared, housing_labels)

tree_reg = DecisionTreeRegressor(random_state=42)
tree_reg.fit(housing_prepared, housing_labels)

forest_reg = RandomForestRegressor(random_state=42)
forest_reg.fit(housing_prepared, housing_labels)
```

➤ Cross Validation RMSE Scores:

```
def display_scores(model, features, labels):
    scores = cross_val_score(model, features, labels,
                              scoring="neg_mean_squared_error", cv=10)
    rmse_scores = np.sqrt(-scores)
    print(f"Model: {model.__class__.__name__}")
    print("Scores:", rmse_scores)
    print("Mean:", rmse_scores.mean())
    print("Standard deviation:", rmse_scores.std())
    print("")
```

➤ Evaluating Models:

```
display_scores(tree_reg, housing_prepared, housing_labels)
display_scores(forest_reg, housing_prepared, housing_labels)
```

➤ Grid Search for best hyperparameters for Random Forest:

```
param_grid = [
    {'n_estimators': [30, 50, 100], 'max_features': [4, 6, 8]},
    {'bootstrap': [False], 'n_estimators': [30, 50], 'max_features': [4, 6]},
]

grid_search = GridSearchCV(forest_reg, param_grid, cv=5,
                           scoring='neg_mean_squared_error',
                           return_train_score=True)
grid_search.fit(housing_prepared, housing_labels)
```

➤ Output best parameter:

```
print("Best parameters from Grid Search:")
print(grid_search.best_params_)
```

➤ Randomized Search:

```
param_dist = {
    'n_estimators': randint(30, 100),
    'max_features': randint(2, 8),
}

random_search = RandomizedSearchCV(forest_reg, param_distributions=param_dist,
                                   n_iter=10, cv=5, scoring='neg_mean_squared_error',
                                   random_state=42, return_train_score=True)
random_search.fit(housing_prepared, housing_labels)

print("Best parameters from Randomized Search:")
print(random_search.best_params_)
```

➤ Evaluate Final Model

```
final_model = random_search.best_estimator_
```

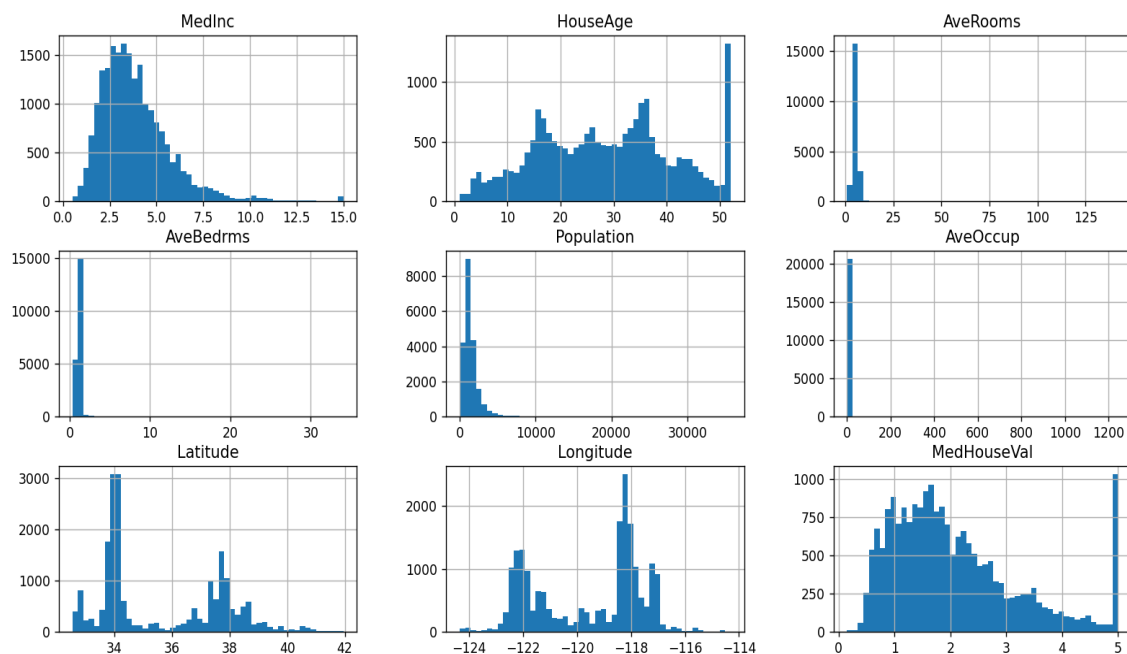
➤ Prepare for the test set:

```
X_test = housing_test.drop("MedHouseVal", axis=1)
y_test = housing_test["MedHouseVal"].copy()
X_test_prepared = full_pipeline.transform(X_test)

final_predictions = final_model.predict(X_test_prepared)
final_mse = mean_squared_error(y_test, final_predictions)
final_rmse = np.sqrt(final_mse)

print("\nFinal Model Evaluation on Test Set:")
print("RMSE:", final_rmse)
```

➤ All Graphs:



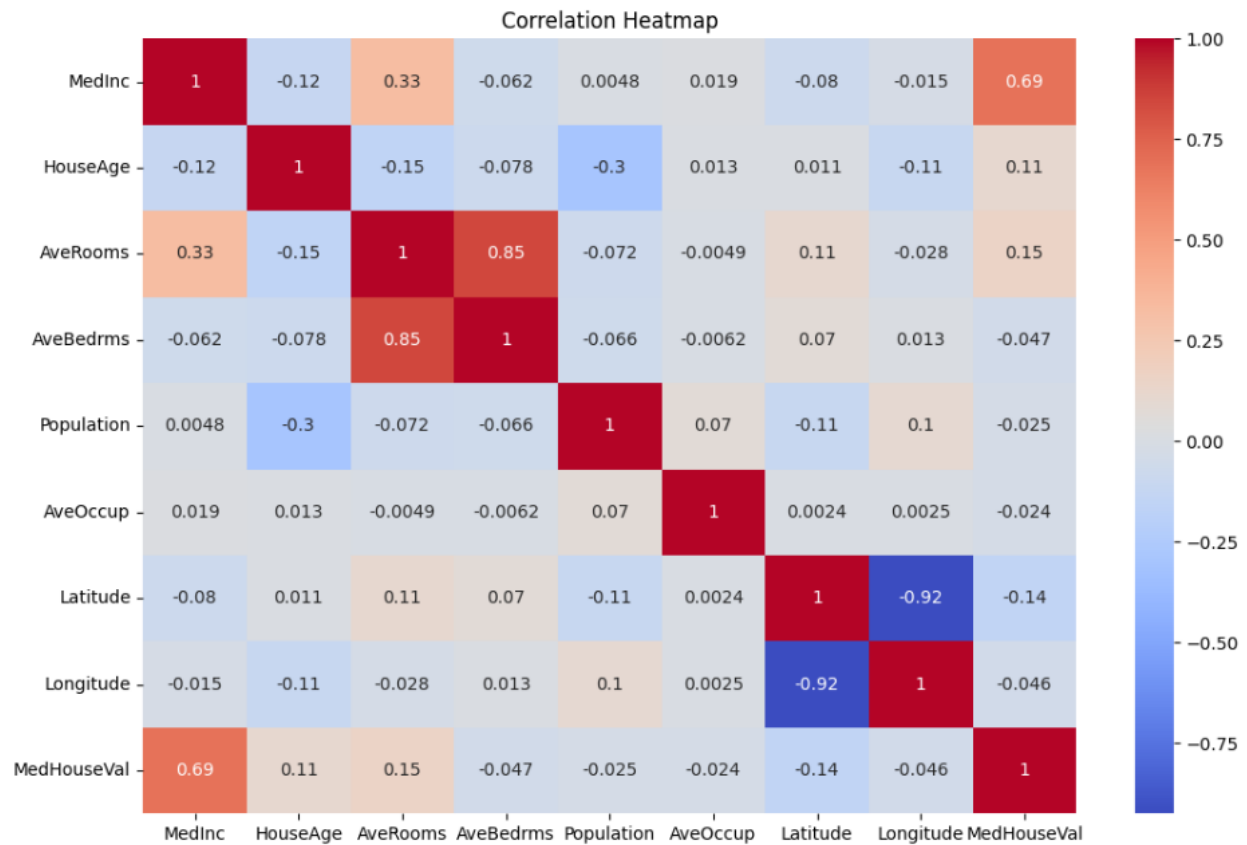
➤ Other Outputs:

| | MedInc | HouseAge | AveRooms | ... | Latitude | Longitude | MedHouseVal |
|---|--------|----------|----------|-----|----------|-----------|-------------|
| 0 | 8.3252 | 41.0 | 6.984127 | ... | 37.88 | -122.23 | 4.526 |
| 1 | 8.3014 | 21.0 | 6.238137 | ... | 37.86 | -122.22 | 3.585 |
| 2 | 7.2574 | 52.0 | 8.288136 | ... | 37.85 | -122.24 | 3.521 |
| 3 | 5.6431 | 52.0 | 5.817352 | ... | 37.85 | -122.25 | 3.413 |
| 4 | 3.8462 | 52.0 | 6.281853 | ... | 37.85 | -122.25 | 3.422 |

```
[5 rows x 9 columns]
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype  
---  -
0   MedInc          20640 non-null  float64
1   HouseAge        20640 non-null  float64
2   AveRooms        20640 non-null  float64
3   AveBedrms       20640 non-null  float64
4   Population      20640 non-null  float64
5   AveOccup        20640 non-null  float64
6   Latitude        20640 non-null  float64
7   Longitude       20640 non-null  float64
8   MedHouseVal     20640 non-null  float64
dtypes: float64(9)
memory usage: 1.4 MB
```

| | MedInc | HouseAge | ... | Longitude | MedHouseVal |
|-------|--------------|--------------|-----|--------------|--------------|
| count | 20640.000000 | 20640.000000 | ... | 20640.000000 | 20640.000000 |
| mean | 3.870671 | 28.639486 | ... | -119.569704 | 2.068558 |
| std | 1.899822 | 12.585558 | ... | 2.003532 | 1.153956 |
| min | 0.499900 | 1.000000 | ... | -124.350000 | 0.149990 |
| 25% | 2.563400 | 18.000000 | ... | -121.800000 | 1.196000 |
| 50% | 3.534800 | 29.000000 | ... | -118.490000 | 1.797000 |
| 75% | 4.743250 | 37.000000 | ... | -118.010000 | 2.647250 |
| max | 15.000100 | 52.000000 | ... | -114.310000 | 5.000010 |

[8 rows x 9 columns]



```
Model: DecisionTreeRegressor
Scores: [0.71833277 0.76445327 0.68953762 0.73944007 0.72894604 0.68827585
 0.71477233 0.744066 0.74345713 0.77495355]
Mean: 0.730623462355596
Standard deviation: 0.02723829991358953
```

```
Model: RandomForestRegressor
Scores: [0.48250658 0.52139465 0.5009443 0.52172245 0.52180332 0.48496938
 0.48770311 0.51557347 0.50358903 0.52014149]
Mean: 0.5060347778945158
Standard deviation: 0.015446467010870067
```

```
Best parameters from Grid Search:
{'bootstrap': False, 'max_features': 4, 'n_estimators': 50}
```

```
Best parameters from Randomized Search:
{'max_features': 4, 'n_estimators': 90}
```

```
Final Model Evaluation on Test Set:
RMSE: 0.5017152121806703
```